

# Spark


Part 3:

*Data frames, SQL.*

Dario Colazzo


# Spark programming model

- ◉ **RDDs** : collection of element values distributed over the cluster, mainly in main-memory (RAM)
- ◉ **Transformations** : *lazy* operators that create *new* RDDs from RDDs.
- ◉ **Actions** : launch a *computation* and return a *value* to the program driver or write data to the *external storage*



# Dataframes

# Dataframes and Spark SQL




# Dataframe

- ◉ A Dataframe is a distributed collection of rows
- ◉ Homogeneous schema
- ◉ Somewhat equivalent to a table in a relational database.

# Adding schema to RDDs

- Spark+RDD: **functional** transformations on partitioned collections of **opaque** objects.
- SQL + DataFrame: **declarative** transformations on partitioned collections of **tuples**.



# Creating DataFrames

```
scala> val spark = SparkSession
        .builder()
        .appName("Spark SQL basic example")
        .config("abc", "cdf")
        .getOrCreate()
```

```
scala> import spark.implicits._
import spark.implicits._
```

```
val df = spark
    .read
    .json("/usr/local/Cellar/apache-spark/2.4.0/libexec/
examples/src/main/resources/people.json")
```



Local path

# Creating DataFrames

```
scala> df.show()
```

```
+-----+-----+  
|  age |   name |  
+-----+-----+  
| null | Michael |  
|   30 |    Andy |  
|   19 |   Justin |  
+-----+-----+
```

```
scala>
```



# Using DataFrames

```
scala> df.select($"name").show()
```

```
+-----+  
|  name |  
+-----+  
|Michael|  
|  Andy |  
| Justin|  
+-----+
```

Select everybody, but  
increment the age by 1

```
scala> df.select($"name", $"age" + 1).show()
```

```
+-----+-----+  
|  name | (age + 1) |  
+-----+-----+  
|Michael|      null |  
|  Andy |       31 |  
| Justin|       20 |  
+-----+-----+
```

# Using DataFrames

```
scala> df.filter($"age" > 21).show()
```

```
+----+-----+  
| age | name |  
+----+-----+  
|  30 | Andy |  
+----+-----+
```

```
scala> df.groupBy($"age").count().show()
```

```
+----+-----+  
| age | count |  
+----+-----+  
|  19 |      1 |  
| null |      1 |  
|  30 |      1 |  
+----+-----+
```

# SQL on DataFrames

Register the DataFrame as a SQL temporary view

```
scala> df.createOrReplaceTempView("people")
```

```
scala> val sqlDF = spark.sql("SELECT * FROM people")
```

```
scala> sqlDF.show()
```

```
+-----+-----+
|  age|   name|
+---+-----+
| null|Michael|
|   30|   Andy|
|   19| Justin|
+-----+-----+
```

```
scala>
```

# Converting RDDs into DataFrames

```
scala> val spark = SparkSession
      .builder()
      .appName("Spark SQL basic example")
      .config("abc", "cdf")
      .getOrCreate()
```

```
scala> import spark.implicits._
import spark.implicits._
```

# Converting RDDs into DataFrames

```
scala> val data = Array(("a",1), ("b",2) , ("a",3) , ("c",4), ("b",5))  
data: Array[(String, Int)] = Array((a,1), (b,2), (a,3), (c,4), (b,5))
```

```
scala> val rdd = sc.parallelize(data)  
rdd: org.apache.spark.rdd.RDD[(String, Int)] =  
ParallelCollectionRDD[47] at parallelize at <console>:35
```

```
scala> val rdd_1 = rdd.reduceByKey((a, b) => a + b)
```

```
scala> rdd_1.collect  
res13: Array[(String, Int)] = Array((a,4), (b,7), (c,4))
```

```
scala>
```

# Converting RDDs into DataFrames

```
scala> rdd_1.collect
res13: Array[(String, Int)] = Array((a,4), (b,7), (c,4))

scala> val myDf = rdd_1.toDF("name", "val")
myDf: org.apache.spark.sql.DataFrame = [name: string, val: int]

scala> myDf.show()
+-----+-----+
| name | val |
+-----+-----+
|    a |   4 |
|    b |   7 |
|    c |   4 |
+-----+-----+

scala> myDf.printSchema
root
 |-- name: string (nullable = true)
 |-- val: integer (nullable = false)
scala>
```

More details in this nice post

<https://indatalabs.com/blog/convert-spark-rdd-to-dataframe-dataset>

# Conclusion

- Supports on a **variety** of data sources.



- A DataFrame can be operated on as **normal RDDs** or as a **temporary** table.
- Registering a DataFrame as a **table** allows you to run SQL queries over its data.
- More details on :

<http://spark.apache.org/docs/latest/sql-programming-guide.html#starting-point-sparksession>

# Datasets



# Datasets

- Datasets offers a compromise/mix between RDD and DataFrames
  - They can be used to run SQL operations
- A Dataset is a collection of JVM objects, so objects are strongly typed

# First example

```
scala> val wordsRDD = sc.parallelize(Seq("Spark I am your father", "May the spark be with you", "Spark I am your father"))
```

```
wordsRDD: org.apache.spark.rdd.RDD[String] = ParallelCollectionRDD[59] at parallelize at <console>:33
```

```
scala> val wordsDataset = wordsRDD.toDS()
```

```
wordsDataset: org.apache.spark.sql.Dataset[String] = [value: string]
```

```
scala> wordsDataset.printSchema
```

```
root
```

```
 |-- value: string (nullable = true)
```

```
scala> wordsDataset.show()
```

```
+-----+
|           value|
+-----+
|Spark I am your f...|
|May the spark be ...|
|Spark I am your f...|
+-----+
```

# First example

```
scala> val groupedDataset = wordsDataset.flatMap(_.toLowerCase.split(" "))  
                                         .filter(_ != "")  
                                         .groupBy("value")
```

```
scala> val countsDataset = groupedDataset.count()
```

```
scala> countsDataset.show()
```

```
scala> groupedDataset.count().show
```

```
+-----+-----+  
| value|count|  
+-----+-----+  
|father|    2|  
|  you|    1|  
| with|    1|  
|   be|    1|  
| your|    2|  
|  may|    1|  
|spark|    3|
```

# RDD & Datasets

## RDDs

```
val lines = sc.textFile("/wikipedia")  
  
val words = lines  
    .flatMap(_.split(" "))  
    .filter(_ != "")
```

## Datasets

```
val lines = sqlContext.read.text("/wikipedia").as[String]  
  
val words = lines  
    .flatMap(_.split(" "))  
    .filter(_ != "")
```

# Word-count

## RDDs

```
val counts = words

    .groupBy(_.toLowerCase)

    .map(w => (w._1, w._2.size))
```

## Datasets

```
val counts = words

    .groupBy(_.toLowerCase)

    .count()
```

# Performances



# Readings

- Interesting post on serialisation/deserialisation
  - <https://blog.xebia.fr/2017/09/27/spark-comprendre-et-corriger-lexception-task-not-serializable/>
- Datasets:
  - <https://spark.apache.org/docs/2.3.0/api/java/index.html?org/apache/spark/sql/Dataset.html>

# FETS in Spark



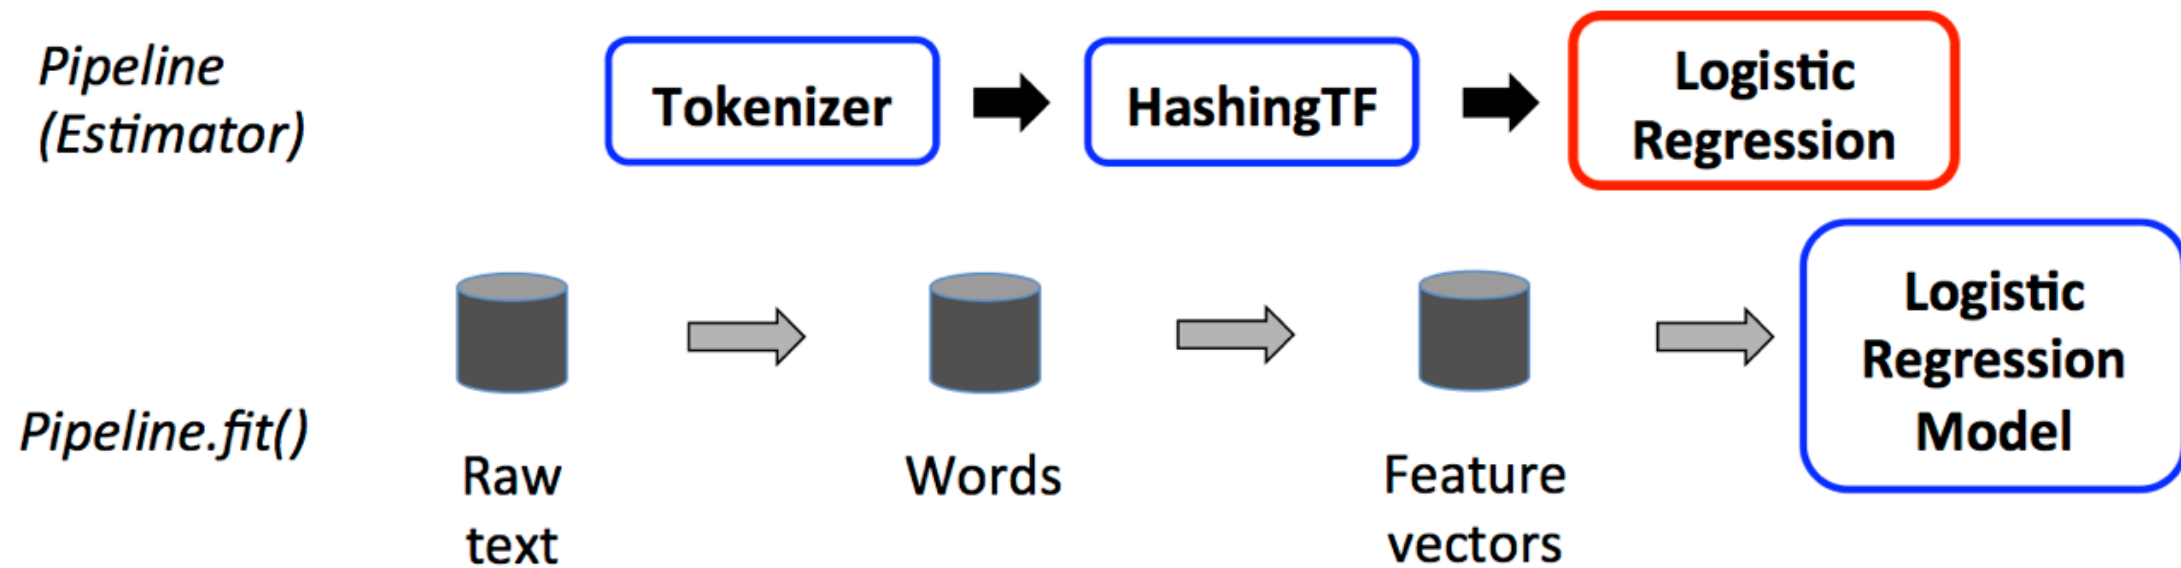
# Spark MLlib

- We will rely on Spark MLlib
- Main notion it relies on: ML pipelines
- Mostly inspired by scikit-learn

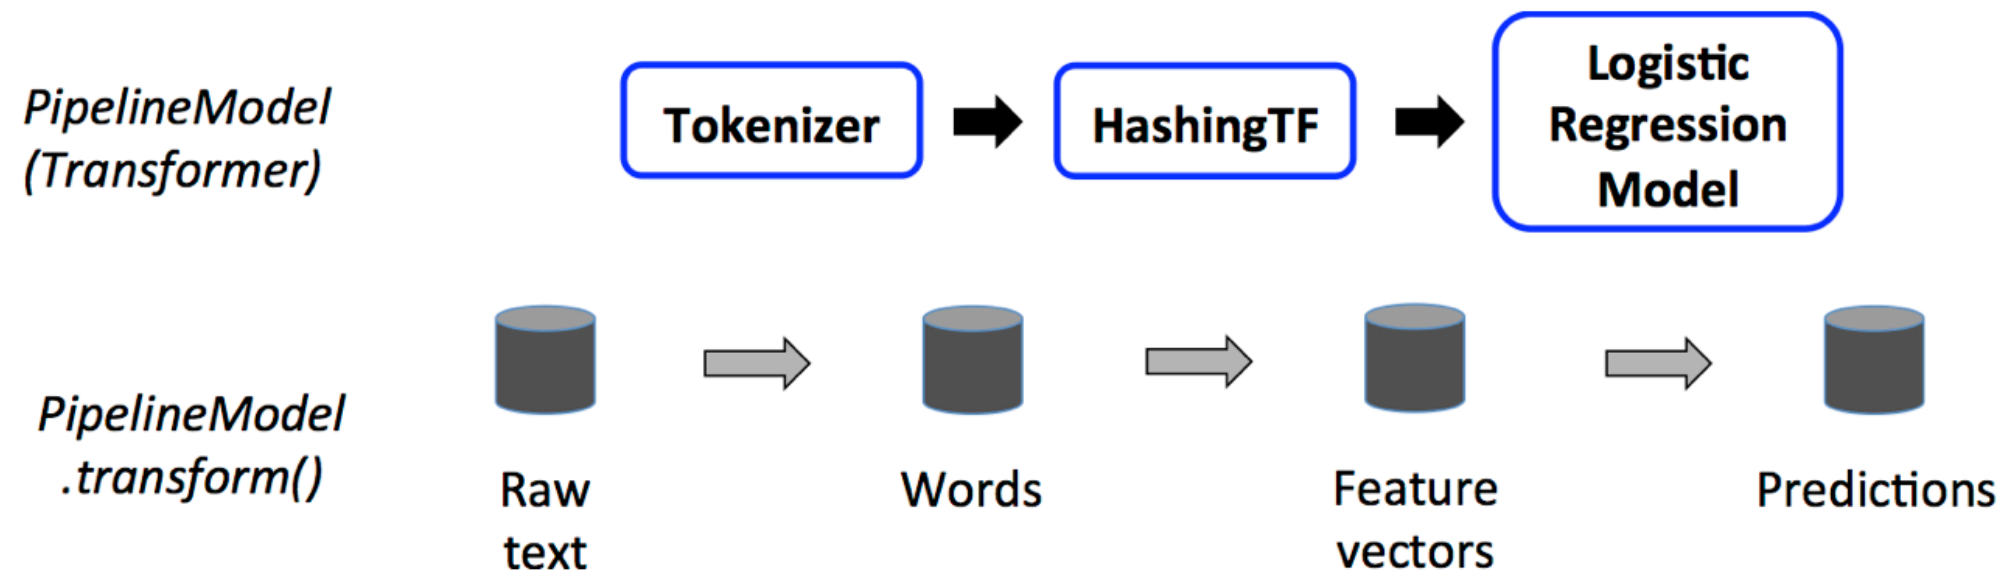
# Pipeline

- A pipeline chains several ML steps/stages in order to specify a ML workflow, including all the steps: preprocessing, feature extraction, model training, estimation, ...
- Each step/stage is either a Transformer or an Estimator
- From Spark documentation:
  - **Transformer**: A Transformer is an algorithm which can transform one DataFrame into another DataFrame. E.g., an ML model is a Transformer which transforms a DataFrame with features into a DataFrame with predictions.
  - **Estimator**: An Estimator is an algorithm which can be fit on a DataFrame to produce a Transformer. E.g., a learning algorithm is an Estimator which trains on a DataFrame and produces a model.

# Training Pipeline



# Pipeline model



# Example, a couple of interesting Transformers

```
scala> import org.apache.spark.ml.feature.{HashingTF, Tokenizer}
```

```
scala> val sentenceData = spark.createDataFrame(Seq(  
  | (0, "Hi I heard about Spark about"),  
  | (0, "I wish Java could case classes"),  
  | (1, "Logistic regression models are neat")  
  | )).toDF("label", "sentence")
```

```
scala> sentenceData.show(3,truncate=false)  
+-----+-----+  
|label|sentence|  
+-----+-----+  
|0     |Hi I heard about Spark|  
|0     |I wish Java could use case classes|  
|1     |Logistic regression models are neat|  
+-----+-----+
```

# HashingTF

```
scala> val tokenizer = new Tokenizer().setInputCol("sentence").setOutputCol("words")
```

```
scala> val hashingTF = new HashingTF().  
    | setInputCol("words").setOutputCol("rawFeatures").setNumFeatures(100)
```

```
scala> val wordsData = tokenizer.transform(sentenceData)
```

```
scala> wordsData.show(1, truncate=false)
```

label	sentence	words
0	Hi I heard about Spark	[hi, i, heard, about, spark]

```
scala> val featurizedData = hashingTF.transform(wordsData)
```

```
scala> featurizedData.show(3, truncate=false)
```

label	sentence	words	rawFeatures
0	Hi I heard about Spark about	[hi, i, heard, about, spark, about]	(100,[56,68,73,86],[3.0,1.0,1.0,1.0])
0	I wish Java could case classes	[i, wish, java, could, case, classes]	(100,[7,42,56,67,80,95],[1.0,1.0,1.0,1.0,1.0,1.0])
1	Logistic regression models are neat	[logistic, regression, models, are, neat]	(100,[4,59,63,71,86],[1.0,1.0,1.0,1.0,1.0])

```
scala> featurizedData.printSchema()
```

```
root
```

```
-- label: integer (nullable = false)  
-- sentence: string (nullable = true)  
-- words: array (nullable = true)  
|   |-- element: string (containsNull = true)  
-- rawFeatures: vector (nullable = true)
```

# CountVectorizer

```
scala> import org.apache.spark.ml.feature.{CountVectorizer,Tokenizer}
```

```
scala> val tokenizer = new Tokenizer().setInputCol("sentence").setOutputCol("words")
```

```
scala> val wordsData = tokenizer.transform(sentenceData)
```

```
scala> wordsData.printSchema()
```

```
root
 |-- label: integer (nullable = false)
 |-- sentence: string (nullable = true)
 |-- words: array (nullable = true)
 |       |-- element: string (containsNull = true)
```

```
scala> val countVectorizer =
      new CountVectorizer().setInputCol("words").setOutputCol("features").setVocabSize(20)
```

```
scala> val featurizedDataModel = countVectorizer.fit(wordsData)
```

```
scala> featurizedDataModel.transform(wordsData).show(3, truncate=false)
```


label	sentence	words	features
0	Hi I heard about Spark	[hi, i, heard, about, spark]	[(16,[0,7,8,9,13],[1.0,1.0,1.0,1.0,1.0])]
0	I wish Java could use case classes	[i, wish, java, could, use, case, classes]	[(16,[0,1,4,5,10,12,15],[1.0,1.0,1.0,1.0,1.0,1.0,1.0])]
1	Logistic regression models are neat	[logistic, regression, models, are, neat]	[(16,[2,3,6,11,14],[1.0,1.0,1.0,1.0,1.0])]

```
scala> featurizedDataModel.vocabulary
```

```
res46: Array[String] = Array(i, could, regression, neat, java, case, models, spark, about, hi, wish, are, classes, heard, logistic, use)
```

# StringIndexer

id	category
0	a
1	b
2	c
3	a
4	a
5	c



id	category	categoryIndex
0	a	0.0
1	b	2.0
2	c	1.0
3	a	0.0
4	a	0.0
5	c	1.0



# StringIndexer

```
scala> import org.apache.spark.ml.feature.StringIndexer

scala> val df = spark.createDataFrame(
  |   Seq((0, "a"), (1, "b"), (2, "c"), (3, "a"), (4, "a"), (5, "c"))
  | ).toDF("id", "category")

scala> val indexer = new StringIndexer().
  | setInputCol("category").
  | setOutputCol("categoryIndex")

scala> val indexed = indexer.fit(df).transform(df)

scala> indexed.show()
+---+-----+-----+
| id|category|categoryIndex|
+---+-----+-----+
|  0|      a|          0.0|
|  1|      b|          2.0|
|  2|      c|          1.0|
|  3|      a|          0.0|
|  4|      a|          0.0|
|  5|      c|          1.0|
+---+-----+-----+
```

# oneHotEncoder

```
scala> import org.apache.spark.ml.feature.OneHotEncoder
```

```
scala> val df = spark.createDataFrame(Seq(  
  |   (0.0, 1.0),  
  |   (1.0, 0.0),  
  |   (2.0, 1.0),  
  |   (0.0, 2.0),  
  |   (0.0, 1.0),  
  |   (2.0, 2.0)  
  | ).toDF("categoryIndex1", "categoryIndex2")
```

```
scala> val encoder = new OneHotEncoder().  
  | setInputCols(Array("categoryIndex1", "categoryIndex2")).  
  | setOutputCols(Array("categoryVec1", "categoryVec2"))
```

```
scala> val model = encoder.fit(df)
```

```
scala> val encoded = model.transform(df).cache()
```

```
scala> encoded.show()
```

categoryIndex1	categoryIndex2	categoryVec1	categoryVec2
0.0	1.0	(2, [0], [1.0])	(2, [1], [1.0])
1.0	0.0	(2, [1], [1.0])	(2, [0], [1.0])
2.0	1.0	(2, [], [])	(2, [1], [1.0])
0.0	2.0	(2, [0], [1.0])	(2, [], [])
0.0	1.0	(2, [0], [1.0])	(2, [1], [1.0])
2.0	2.0	(2, [], [])	(2, [], [])

```
scala>
```

```
scala> encoded.printSchema()  
root  
|-- categoryIndex1: double (nullable = false)  
|-- categoryIndex2: double (nullable = false)  
|-- categoryVec1: vector (nullable = true)  
|-- categoryVec2: vector (nullable = true)
```

```
scala>
```

# VectorAssembler

```
scala> encoded.printSchema()
```

```
root
```

```
|-- categoryIndex1: double (nullable = false)
|-- categoryIndex2: double (nullable = false)
|-- categoryVec1: vector (nullable = true)
|-- categoryVec2: vector (nullable = true)
```

```
scala> import org.apache.spark.ml.feature.VectorAssembler
```

```
scala> val assembler = new VectorAssembler().
  | setInputCols(Array("categoryVec1", "categoryVec2")).
  | setOutputCol("features")
```

```
scala> val output = assembler.transform(encoded)
```

```
scala> output.show()
```

categoryIndex1	categoryIndex2	categoryVec1	categoryVec2	features
0.0	1.0	(2, [0], [1.0])	(2, [1], [1.0])	[1.0, 0.0, 0.0, 1.0]
1.0	0.0	(2, [1], [1.0])	(2, [0], [1.0])	[0.0, 1.0, 1.0, 0.0]
2.0	1.0	(2, [], [])	(2, [1], [1.0])	(4, [3], [1.0])
0.0	2.0	(2, [0], [1.0])	(2, [], [])	(4, [0], [1.0])
0.0	1.0	(2, [0], [1.0])	(2, [1], [1.0])	[1.0, 0.0, 0.0, 1.0]
2.0	2.0	(2, [], [])	(2, [], [])	(4, [], [])

# VectorAssembler - how to avoid information loss

```
scala> encoded.printSchema()
```

```
root
|-- categoryIndex1: double (nullable = false)
|-- categoryIndex2: double (nullable = false)
|-- categoryVec1: vector (nullable = true)
|-- categoryVec2: vector (nullable = true)
```

```
scala> import org.apache.spark.ml.feature.VectorAssembler
```

```
scala> val encoder = new OneHotEncoderEstimator().
      | setInputCols(Array("categoryIndex1", "categoryIndex2")).
      | setOutputCols(Array("categoryVec1", "categoryVec2")).setDropLast(false)
```

```
scala> val model = encoder.fit(df)
model: org.apache.spark.ml.feature.OneHotEncoderModel = oneHotEncoder_9715ada877c5
```

```
scala> val encoded = model.transform(df)
encoded: org.apache.spark.sql.DataFrame = [categoryIndex1: double, categoryIndex2: double ... 2 more fields]
```

```
scala> encoded.show()
```

categoryIndex1	categoryIndex2	categoryVec1	categoryVec2
0.0	1.0	(3, [0], [1.0])	(3, [1], [1.0])
1.0	0.0	(3, [1], [1.0])	(3, [0], [1.0])
2.0	1.0	(3, [2], [1.0])	(3, [1], [1.0])
0.0	2.0	(3, [0], [1.0])	(3, [2], [1.0])
0.0	1.0	(3, [0], [1.0])	(3, [1], [1.0])
2.0	2.0	(3, [2], [1.0])	(3, [2], [1.0])

```
scala>
```

# VectorAssembler - how to avoid information loss

Same as  
before

```
scala> val output = assembler.transform(encoded)
```

```
scala> output.show()
```

categoryIndex1	categoryIndex2	categoryVec1	categoryVec2	features
0.0	1.0	(3, [0], [1.0])	(3, [1], [1.0])	(6, [0,4], [1.0,1.0])
1.0	0.0	(3, [1], [1.0])	(3, [0], [1.0])	(6, [1,3], [1.0,1.0])
2.0	1.0	(3, [2], [1.0])	(3, [1], [1.0])	(6, [2,4], [1.0,1.0])
0.0	2.0	(3, [0], [1.0])	(3, [2], [1.0])	(6, [0,5], [1.0,1.0])
0.0	1.0	(3, [0], [1.0])	(3, [1], [1.0])	(6, [0,4], [1.0,1.0])
2.0	2.0	(3, [2], [1.0])	(3, [2], [1.0])	(6, [2,5], [1.0,1.0])

```
scala>
```

# Let's now obtain an RDD with arrays corresponding to assembled vectors

```
scala> output.show()
```

categoryIndex1	categoryIndex2	categoryVec1	categoryVec2	features
0.0	1.0	(3, [0], [1.0])	(3, [1], [1.0])	(6, [0,4], [1.0,1.0])
1.0	0.0	(3, [1], [1.0])	(3, [0], [1.0])	(6, [1,3], [1.0,1.0])
2.0	1.0	(3, [2], [1.0])	(3, [1], [1.0])	(6, [2,4], [1.0,1.0])
0.0	2.0	(3, [0], [1.0])	(3, [2], [1.0])	(6, [0,5], [1.0,1.0])
0.0	1.0	(3, [0], [1.0])	(3, [1], [1.0])	(6, [0,4], [1.0,1.0])
2.0	2.0	(3, [2], [1.0])	(3, [2], [1.0])	(6, [2,5], [1.0,1.0])

Attention: if you have both Sparse and Dense vectors, just use Vector instead of SparseVecor

```
scala> import org.apache.spark.ml.linalg.SparseVector
```

```
scala> val toArr: Any => Array[Double] = _.asInstanceOf[SparseVector].toArray
```

```
scala> val toArrUdf = udf(toArr)
```

```
scala> val outputWithArrayFeat = output.withColumn("features_arr",toArrUdf('features)).cache()
```

```
scala> outputWithArrayFeat.show(truncate=false)
```

categoryIndex1	categoryIndex2	categoryVec1	categoryVec2	features	features_arr
0.0	1.0	(3, [0], [1.0])	(3, [1], [1.0])	(6, [0,4], [1.0,1.0])	[1.0, 0.0, 0.0, 0.0, 1.0, 0.0]
1.0	0.0	(3, [1], [1.0])	(3, [0], [1.0])	(6, [1,3], [1.0,1.0])	[0.0, 1.0, 0.0, 1.0, 0.0, 0.0]
2.0	1.0	(3, [2], [1.0])	(3, [1], [1.0])	(6, [2,4], [1.0,1.0])	[0.0, 0.0, 1.0, 0.0, 1.0, 0.0]
0.0	2.0	(3, [0], [1.0])	(3, [2], [1.0])	(6, [0,5], [1.0,1.0])	[1.0, 0.0, 0.0, 0.0, 0.0, 1.0]
0.0	1.0	(3, [0], [1.0])	(3, [1], [1.0])	(6, [0,4], [1.0,1.0])	[1.0, 0.0, 0.0, 0.0, 1.0, 0.0]
2.0	2.0	(3, [2], [1.0])	(3, [2], [1.0])	(6, [2,5], [1.0,1.0])	[0.0, 0.0, 1.0, 0.0, 0.0, 1.0]

# Let's now obtain an RDD with arrays corresponding to assembled vectors

```
scala> val outputWithArrayOnlyFeat = outputWithArrayFeat.select("features_arr")
```

```
scala> outputWithArrayOnlyFeat.rdd.take(1)
res11: Array[org.apache.spark.sql.Row] = Array([WrappedArray(1.0, 0.0, 0.0, 0.0, 1.0, 0.0)])
```

```
scala> outputWithArrayOnlyFeat.rdd.map(x => x(0))
res18: org.apache.spark.rdd.RDD[Any] = MapPartitionsRDD[62] at map at <console>:32
```

```
scala> import scala.collection.mutable.WrappedArray
```

```
scala> val myrdd = outputWithArrayOnlyFeat.rdd.map(x => x(0).asInstanceOf[WrappedArray[Double]].toArray[Double])
myrdd: org.apache.spark.rdd.RDD[Array[Double]] = MapPartitionsRDD[64] at map at <console>:31
```

```
scala> myrdd.take(1)
res20: Array[Array[Double]] = Array(Array(1.0, 0.0, 0.0, 0.0, 1.0, 0.0))
```

# Pipelines

- Once data preparation steps and model estimator are ready you can concatenate them in a pipeline
- For instance, a pipeline for data preparation can be built and used as follows



# Pipeline

```
scala> import org.apache.spark.ml.feature.OneHotEncoder
```

```
scala> import org.apache.spark.ml.feature.VectorAssembler
```

```
scala> val df = spark.createDataFrame(Seq(  
  |   (0.0, 1.0),  
  |   (1.0, 0.0),  
  |   (2.0, 1.0),  
  |   (0.0, 2.0),  
  |   (0.0, 1.0),  
  |   (2.0, 2.0)  
  |   )).toDF("categoryIndex1", "categoryIndex2")
```

```
scala> val encoder = new OneHotEncoder().  
  | setInputCols(Array("categoryIndex1", "categoryIndex2")).  
  | setOutputCols(Array("categoryVec1", "categoryVec2"))
```

```
scala> val assembler = new VectorAssembler().  
  | setInputCols(Array("categoryVec1", "categoryVec2")).  
  | setOutputCol("features")
```

```
scala> val steps: Array[org.apache.spark.ml.PipelineStage] = Array(encoder, assembler)
```

```
scala> import org.apache.spark.ml.Pipeline
```

```
scala> val pipeline_prep = new Pipeline().setStages(steps)
```

```
scala> val output = (pipeline_prep.fit(df)).transform(df)
```

```
scala> output.show(truncate=false)
```

categoryIndex1	categoryIndex2	categoryVec1	categoryVec2	features
0.0	1.0	(2, [0], [1.0])	(2, [1], [1.0])	[1.0, 0.0, 0.0, 1.0]
1.0	0.0	(2, [1], [1.0])	(2, [0], [1.0])	[0.0, 1.0, 1.0, 0.0]
2.0	1.0	(2, [], [])	(2, [1], [1.0])	(4, [3], [1.0])
0.0	2.0	(2, [0], [1.0])	(2, [], [])	(4, [0], [1.0])
0.0	1.0	(2, [0], [1.0])	(2, [1], [1.0])	[1.0, 0.0, 0.0, 1.0]
2.0	2.0	(2, [], [])	(2, [], [])	(4, [], [])

# Data splitting

- In order to split a dataset for training and testing you can proceed as follows

```
scala> val Array(training, test) = data.randomSplit(Array(0.75, 0.25), seed = 12345 )
```

# Building a model estimator and parameters

```
scala> import org.apache.spark.ml.regression.{LinearRegression}
```

```
scala> import org.apache.spark.ml.tuning.{ParamGridBuilder, TrainValidationSplit}
```

```
scala> val lr = new LinearRegression().setLabelCol("priceOutputVar").setFeaturesCol("features")
```

```
scala> val paramGrid = new ParamGridBuilder().  
    | addGrid(lr.regParam, Array(0.1, 0.01)).  
    | addGrid(lr.fitIntercept).  
    | addGrid(lr.elasticNetParam, Array(0.0, 1.0)).build()
```

```
paramGrid: Array[org.apache.spark.ml.param.ParamMap] =  
Array({
```

```
  linReg_41a489042a50-elasticNetParam: 0.0,  
  linReg_41a489042a50-fitIntercept: true,  
  linReg_41a489042a50-regParam: 0.1
```

```
}, {  
  linReg_41a489042a50-elasticNetParam: 0.0,  
  linReg_41a489042a50-fitIntercept: true,  
  linReg_41a489042a50-regParam: 0.01
```

```
}, {  
  linReg_41a489042a50-elasticNetParam: 1.0,  
  linReg_41a489042a50-fitIntercept: true,  
  linReg_41a489042a50-regParam: 0.1
```

```
}, {  
  linReg_41a489042a50-elasticNetParam: 1.0,  
  linReg_41a489042a50-fitIntercept: true,  
  linReg_41a489042a50-regParam: 0.01
```

```
}, {  
  linReg_41a489042a50-elasticNetParam: 0.0,  
  linReg_41a489042a50-fitIntercept: false,  
  linReg_41a489042a50-regParam: 0.1
```

```
}, {  
  linReg_41a489042a50-elasticNetParam: 0.0,  
  linReg_41a489042a50-fitIntercept: false,  
  linReg_41a489042a50-regPa...
```

```
scala>
```

Features and label  
to be estimated  
have been already  
prepared

$$\alpha (\lambda \|\mathbf{w}\|_1) + (1 - \alpha) \left( \frac{\lambda}{2} \|\mathbf{w}\|_2^2 \right), \alpha \in [0, 1], \lambda \geq 0$$

**Elastic net regularization:**  
you have L1 for  $\alpha=1$  while  
L2 is obtained for  $\alpha=0$ .  
Lambda is the reg-  
param.

# A complete pipeline (to be used for Diamonds regression )

```
scala> val steps: Array[org.apache.spark.ml.PipelineStage] =  
      categoricalIndexers ++ categoricalEncoders ++ Array(assembler, lr)  
  
scala> val pipeline = new Pipeline().setStages(steps)  
  
scala> import org.apache.spark.ml.evaluation.{RegressionEvaluator}  
  
scala> val tvs = new TrainValidationSplit().  
      | setEstimator(pipeline).  
      | setEvaluator( new RegressionEvaluator().setLabelCol("priceOutputVar") ).  
      | setEstimatorParamMaps(paramGrid).  
      | setTrainRatio(0.75)  
  
scala> val Array(training, test) = data.randomSplit(Array(0.75, 0.25), seed = 12345)  
  
scala> val model = tvs.fit(training)  
  
scala> val holdout = model.transform(test).select("prediction", "priceOutputVar")
```

# Regression model evaluation

Regression analysis is used when predicting a continuous output variable from a number of independent variables.

## Available metrics

Metric	Definition
Mean Squared Error (MSE)	$MSE = \frac{\sum_{i=0}^{N-1} (y_i - \hat{y}_i)^2}{N}$
Root Mean Squared Error (RMSE)	$RMSE = \sqrt{\frac{\sum_{i=0}^{N-1} (y_i - \hat{y}_i)^2}{N}}$
Mean Absolute Error (MAE)	$MAE = \sum_{i=0}^{N-1}  y_i - \hat{y}_i $
Coefficient of Determination ( $R^2$ )	$R^2 = 1 - \frac{MSE}{\text{VAR}(y) \cdot (N-1)} = 1 - \frac{\sum_{i=0}^{N-1} (y_i - \hat{y}_i)^2}{\sum_{i=0}^{N-1} (y_i - \bar{y})^2}$
Explained Variance	$1 - \frac{\text{VAR}(y - \hat{y})}{\text{VAR}(y)}$

<https://spark.apache.org/docs/2.2.0/mllib-evaluation-metrics.html>

# Regression metrics

```
scala> val tvs = new TrainValidationSplit().
  | setEstimator(pipeline).
  | setEvaluator( new RegressionEvaluator().setLabelCol("priceOutputVar") ).
  | setEstimatorParamMaps(paramGrid).
  | setTrainRatio(0.75)

scala> val Array(training, test) = data.randomSplit(Array(0.75, 0.25), seed = 12345)

scala> val model = tvs.fit(training)

scala> val holdout = model.transform(test).select("prediction", "priceOutputVar")

scala> import org.apache.spark.mllib.evaluation.RegressionMetrics

scala> val rm = new RegressionMetrics(holdout.rdd.map(x =>(x(0).asInstanceOf[Double], x(1).asInstanceOf[Double])))
scala> println("sqrt(MSE): " + Math.sqrt(rm.meanSquaredError))
sqrt(MSE): 1141.7457210656032

scala> println("R Squared: " + rm.r2)
R Squared: 0.9179908957599221


scala> println("Explained Variance: " + rm.explainedVariance + "\n")
Explained Variance: 1.4616693408985674E7
```

# What else?

Many other operations:

- <https://spark.apache.org/docs/latest/ml-features.html>

# Diamonds





# FETS for Diamonds

- Dataset available here
  - <https://www.dropbox.com/s/t6jxtilvn31e2gb/diamonds.csv>
- Use spark-shell
- Loading the data:

```
val data = spark.read  
  |.format("csv").option("header", "true").option("inferSchema", true)  
  |.load(datapath).withColumn("priceOutputVar", $"price".cast("double")).cache()
```